

# Variational Autoencoders

①

Let  $x$  denote the vector of all observed variables whose joint distribution  $p^*(x)$  we would like to model. We will attempt to approximate this underlying distribution with a model  $p_\theta(x)$  with parameters

$$p_\theta(x) \approx p^*(x)$$

We will include latent variables  $z$  in our model. These variables are not observed and are not part of the data, but they participate in the generative process producing the observations  $x$ . Now we have a joint distribution  $p_\theta(x, z)$  and the marginal over the observed variables is

$$p_\theta(x) = \int p_\theta(x, z) dz \quad (*)$$

Commonly in the factorization

$$p_\theta(x, z) = p_\theta(z) p_\theta(x|z),$$

The prior distribution  $p_\theta(z)$  and/or  $p_\theta(x|z)$  are specified.

The main difficulty here is that the integral in  $(*)$  is intractable. This also makes the posterior distribution

$$p_\theta(z|x) = \frac{p_\theta(x, z)}{p_\theta(x)}$$

intractable.

To tackle this issue let's introduce an **encoder (inference model)**  $q_\phi(z|x)$ . It has parameters  $\phi$  which we have to optimize so that

$$q_\phi(z|x) \approx p_\theta(z|x)$$

Ex: The distribution  $q_\phi$  can be parametrized using NN's. In that case  $\phi$  includes the weights and the biases of the NN. For example,

$$(\mu, \log \sigma) = \text{Encoder NN}_\phi(x)$$

$$q_\phi(z|x) = \mathcal{N}(z; \mu, \text{diag}(\sigma))$$

# Variational Autoencoders

①

Notice that:

$$\begin{aligned}\log p_{\theta}(x) &= \mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x)] = \mathbb{E}_{q_{\phi}(z|x)} \left[ \log \left[ \frac{p_{\theta}(x, z)}{p_{\theta}(z|x)} \right] \right] = \\ &= \underbrace{\mathbb{E}_{q_{\phi}(z|x)} \left[ \log \left[ \frac{p_{\theta}(x, z)}{q_{\phi}(z|x)} \right] \right]}_{\mathcal{L}_{\theta, \phi}(x)} + \underbrace{\mathbb{E}_{q_{\phi}(z|x)} \left[ \log \left[ \frac{q_{\phi}(z|x)}{p_{\theta}(z|x)} \right] \right]}_{D_{KL}(q_{\phi}(z|x) \| p_{\theta}(z|x))}\end{aligned}$$

The second term above is the Kullback-Liebler (KL) divergence between  $q_{\phi}(z|x)$  and  $p_{\theta}(z|x)$  and is non-negative:

$$D_{KL}(q_{\phi}(z|x) \| p_{\theta}(z|x)) \geq 0$$

The first term  $\mathcal{L}_{\theta, \phi}(x)$  is called **ELBO (evidence lower bound)**.

$$\mathcal{L}_{\theta, \phi}(x) = \mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x, z) - \log q_{\phi}(z|x)] \leq \log p_{\theta}(x) \quad (**)$$

↑  
since  $D_{KL} \geq 0$ .

Note that the KL divergence  $D_{KL}(q_{\phi}(z|x) \| p_{\theta}(z|x))$  determines two distances:

- 1) The distance between the approximate and true posteriors.
- 2) The gap between ELBO  $\mathcal{L}_{\theta, \phi}(x)$  and the marginal likelihood  $\log p_{\theta}(x)$ .

**Two for one:** The maximization of the ELBO  $\mathcal{L}_{\theta, \phi}(x)$  w.r.t. the parameters  $\theta$  and  $\phi$ , will concurrently optimize two desirable objectives:

- 1) It will (implicitly) maximize the marginal likelihood  $p_{\theta}(x) \Rightarrow$

The generative model will become better.

2) It will minimize the distance between the approximation  $q_{\phi}(z|x)$  and the "true" posterior  $p_{\theta}(z|x)$ . The inference model will become better.

Gradients of the ELBO wrt the generative model parameters are easy to obtain:

$$\begin{aligned}\nabla_{\theta} \mathcal{L}_{\theta, \phi}(x) &= \nabla_{\theta} \mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x, z) - \log q_{\phi}(z|x)] = \\ &= \mathbb{E}_{q_{\phi}(z|x)} [\nabla_{\theta} \log p_{\theta}(x, z)] \approx \nabla_{\theta} \log p_{\theta}(x, z)\end{aligned}$$

MC estimator;  $z \sim p_{\theta}(z|x) \leftarrow$  random sample

## Variational Autoencoders

(2)

Gradients wrt the variational parameters  $\phi$  are more difficult

$$\nabla_{\phi} \mathcal{L}_{\theta, \phi}(x) = \nabla_{\phi} \mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x, z) - \log q_{\phi}(z|x)]$$

cannot push the gradient through the expectation.

The reparametrization trick:

Let's express the RV  $z \sim q_{\phi}(z|x)$  as a smooth, bijective transformation of another RV  $u$ , given  $x$  and  $\phi$ :

$$z = g(u, \phi, x),$$

where the distribution of the RV  $u$  is independent of  $x$  or  $\phi$ . The expectation can be rewritten in terms of  $u$ :

$$\mathbb{E}_{q_{\phi}(z|x)} [f(z)] = \mathbb{E}_{p(u)} [f(g(u, \phi, x))]$$

Now the gradient operator and the expectation commute:

$$\begin{aligned} \nabla_{\phi} \mathbb{E}_{q_{\phi}(z|x)} [f(z)] &= \nabla_{\phi} \mathbb{E}_{p(u)} [f(g(u, \phi, x))] = \\ &= \mathbb{E}_{p(u)} [\nabla_{\phi} f(g(u, x, \phi))] \end{aligned}$$

MC estimator;  $u \sim p(u) \leftarrow$  random sample

### Gradient of ELBO.

The ELBO can be rewritten as:

$$\mathcal{L}_{\theta, \phi}(x) = \mathbb{E}_{p(u)} [\log p_{\theta}(x, z) - \log q_{\phi}(z|x)], \quad z = g(u, \phi, x)$$

We can form a simple MC estimator  $\tilde{\mathcal{L}}_{\theta, \phi}(x)$  of the individual-datapoint ELBO where we use a single noise sample  $u \sim p(u)$

$$\tilde{\mathcal{L}}_{\theta, \phi}(x) = \log p_{\theta}(x, z) - \log q_{\phi}(z|x), \quad z = g(u, \phi, x).$$

We can then optimize ELBO using minibatch SGD. Notice also that the gradient is an unbiased estimator of the exact single-datapoint ELBO gradient; when averaged over the noise  $u$ , the estimated gradient matches the 'true' gradient:

$$\mathbb{E}_{p(u)} [\nabla_{\theta, \phi} \tilde{\mathcal{L}}_{\theta, \phi}(x; u)] = \nabla_{\theta, \phi} \mathcal{L}_{\theta, \phi}(x).$$



# Variational Autoencoders

3

## Computation of $\log q_\phi(z|x)$

We know the density  $p(u)$  of the 'noise' distribution. And we have

$$\log q_\phi(z|x) = \log p(u) - \log \left| \det \left( \frac{\partial z}{\partial u} \right) \right|$$

Jacobian matrix

Ex: For example, we can choose

$$u \sim N(0, I)$$

$$(\mu, \log \sigma) = \text{EncoderNN}_\phi(x)$$

$$z = \mu + \sigma \odot u \quad (= g(u, \phi, x)).$$

The Jacobian of this transformation from  $u$  to  $z$  is:

$$\frac{\partial z}{\partial u} = \text{diag}(\sigma)$$

and then

$$\log \left| \det \left( \frac{\partial z}{\partial u} \right) \right| = \sum_i \log \sigma_i$$

and the posterior density is

$$\begin{aligned} \log q_\phi(z|x) &= \log p(u) - \log \left| \det \left( \frac{\partial z}{\partial u} \right) \right| = \\ &= \sum_i \log N(u_i; 0, 1) - \log \sigma_i = -\sum_i \left( \frac{1}{2} (u_i^2 + \log(2\pi)) + \log \sigma_i \right). \end{aligned}$$

Ex: Full covariance Gaussian posterior.

$$u \sim N(0, I)$$

$$z = \mu + Lu,$$

where  $L$  is a lower (or upper) triangular matrix with nonzero entries on the diagonal. Now

$$\log q_\phi(z|x) = \log p(u) - \sum_i \log |L_{ii}|$$

Note also that the covariance of  $z$  is  $\Sigma = LL^T$ . As usual the parameters are computed with a neural network:

$$(\mu, L) \leftarrow \text{EncoderNN}_\phi(x)$$

## Computation of $\log p_\theta(x, z)$ .

Notice that

$$\log p_\theta(x, z) = \log p_\theta(x|z) + \log p_\theta(z)$$

Since  $p_\theta(z) \sim \mathcal{N}(0, I)$  we have

$$\log p_\theta(z) = -\sum_i \frac{1}{2} (z_i^2 + \log(2\pi))$$

For the decoder we assume a factorized Bernoulli distribution

$$p = \text{Decoder NN}_\theta(z)$$

$$\begin{aligned} \log p(x|z) &= \sum_{j=1}^D \log p(x_j|z) = \sum_{j=1}^D \log(\text{Bernoulli}(x_j; p_j)) \\ &= \sum_{j=1}^D x_j \log p_j + (1-x_j) \log(1-p_j) \end{aligned}$$

**Algorithm:** Computation of an unbiased estimate of single-datapoint ELBO for VAE with full-covariance Gaussian inference model and a factorized Bernoulli generative model:

$x$  - data point;  $u$  - random sample from  $p(u) \sim \mathcal{N}(0, I)$ ;

$\theta$  - generative model params;  $\phi$  - inference model params

$q_\phi(z|x)$  - inference model;  $p_\theta(x, z)$  - generative model

$\tilde{\mathcal{L}}$  - unbiased estimate of the single-datapoint ELBO  $\mathcal{L}_{\theta, \phi}(x)$

$$(\mu, L) \leftarrow \text{Encoder NN}_\phi(x)$$

$$u \sim \mathcal{N}(0, I)$$

$$z \leftarrow L u + \mu$$

$$\mathcal{L}^1 = \tilde{\mathcal{L}}[\log q_\phi(z|x)] \leftarrow -\sum_i \left( \frac{1}{2} (u_i^2 + \log(2\pi)) + \log \sigma_i \right)$$

$$\mathcal{L}^2 = \tilde{\mathcal{L}}[\log p_\theta(z)] \leftarrow -\sum_i \left( \frac{1}{2} (z_i^2 + \log(2\pi)) \right)$$

$$p \leftarrow \text{Decoder NN}_\theta(z)$$

$$\mathcal{L}^3 = \tilde{\mathcal{L}}[\log p_\theta(x|z)] \leftarrow \sum_i (x_i \log p_i + (1-x_i) \log(1-p_i)) \quad \tilde{\mathcal{L}} = -\mathcal{L}^1 + \mathcal{L}^2 + \mathcal{L}^3$$